Multi-Object Tracking for Sports Analytics: Player and Ball Tracking with YOLO11n, ByteTrack, a Kalman Filter and Motion-Based Ball Flights

Ravva Siddhartha
Dept. of Computer Science and Engineering
Amrita School of Computing
Amrita Vishwa Vidyapeetham
Bengaluru, India
BL.EN.U4CSE23246

Pulapa Yaswanth Kumar
Dept. of Computer Science and Engineering
Amrita School of Computing
Amrita Vishwa Vidyapeetham
Bengaluru, India
BL.EN.U4CSE23245

Dr. Tripty Singh
Dept. of Computer Science and Engineering
Amrita School of Computing
Amrita Vishwa Vidyapeetham
Bengaluru, India

***Abstract*—This paper presents a lightweight tracking-by-detection system that tracks soccer players and the ball in monocular video on a CPU. A COCO-pretrained YOLO11n detector finds people and the ball in a single pass per frame, ByteTrack assigns persistent player identities, and a constant-velocity Kalman filter tracks the ball and predicts its position through short detection gaps, labelling every predicted position as such. When the ball is struck and the detector loses the motion-blurred ball, camera-motion-compensated moving blobs linked into fast, smooth flights that must start at the ball let the tracker follow it. The proposed player tracker is compared with SORT and OC-SORT on identical cached detections from six test sequences of SportsMOT and UVY-Track (2,950 frames), with hyper-parameters chosen on a separate sequence. The ball tracker is compared with detection-only output on 13 sequences of UVY-Track and InFactory (2,402 frames with a visible ball). On SportsMOT, ByteTrack gives the fewest identity switches (51, against 60 for SORT and 106 for OC-SORT) and the highest recall (95.9%), while tuned SORT reaches a higher HOTA (65.9 against 62.6) because it outputs fewer boxes on people the benchmark does not annotate. On the partially labelled UVY-Track clips ByteTrack has the best OWTA (56.8). The proposed ball tracker raises ball precision from 66.4% to 70.0%, recall from 34.8% to 36.4% and F1 from 45.7% to 47.9% over detection alone and merges 175 correct track segments into 44; on a free kick its motion flights raise ball recall from 91.1% to 99.4%. The detector accounts for 64% of the processing time and is the main limit on accuracy. The proposed method is therefore not uniformly better than the baselines; the paper reports where it helps, where it does not, and why.**

***Index Terms*—multi-object tracking, sports analytics, soccer, ball tracking, YOLO, ByteTrack, SORT, OC-SORT, Kalman filter, HOTA, frame differencing**

I. Introduction

Broadcast and amateur sports video is increasingly analysed automatically for coaching, scouting, officiating support and fan engagement. Almost every such analysis (distance covered, possession, heatmaps, pass maps, tactical shape) starts from the same two signals: where every player is over time, with a stable identity, and where the ball is. Producing these signals from a single camera is a multi-object tracking (MOT) problem with sport-specific difficulties: players of one team look alike, so appearance gives little help in telling them apart; motion is fast and non-linear, with frequent occlusion in duels and set pieces; broadcast cameras pan and zoom, so image motion mixes camera and player motion; and the ball is tiny (often 5–20 pixels), blurred, frequently hidden, and moves much faster than any player.

This work builds a complete, lightweight tracking-by-detection system for this task and evaluates it quantitatively. It runs on a CPU, uses only pretrained models, and is organised so that each stage (detector, player tracker, ball tracker) can be replaced and measured independently. The contributions are:

- a working pipeline that turns a soccer video into an annotated video, a per-frame tracking table and a run summary, in which ball positions predicted by a motion model are always distinguished from detected ones;
- a controlled comparison of ByteTrack with SORT and OC-SORT in which all trackers replay identical detections and are tuned with the same budget on a sequence disjoint from the test data;
- an evaluation on three public datasets with different annotation schemes, including an analysis of how partially labelled data invalidates false-positive-based metrics;
- a ball-tracking evaluation that reports the benefit and the cost of motion prediction separately;
- a motion-based extension of the ball tracker that follows a struck, motion-blurred ball which the generic detector misses, by linking camera-motion-compensated moving blobs into physically plausible flights that must start at the ball.

II. Related Work

Tracking by detection. Modern MOT systems detect objects in every frame and associate detections over time. SORT [1] showed that a Kalman filter [9] for motion prediction plus assignment on bounding-box overlap is a strong and very fast baseline. ByteTrack [2] observed that discarding low-confidence detections throws away true objects that are occluded or blurred, and introduced a second association stage that matches the remaining tracks to low-confidence boxes. OC-SORT [3] kept the simplicity of SORT but corrected the error accumulated by the Kalman filter during occlusion with an observation-centric re-update.

Tracking in sports. SportsMOT [6] is a main public benchmark for player tracking; it emphasises fast, variable-speed motion and similar appearance and identifies association, rather than detection, as the key difficulty in sports. UVY-Track [7] targets unconstrained user-generated sports video rather than broadcast footage, and the InFactory dataset [8] labels the ball in professional soccer broadcasts.

Ball detection and tracking. The ball is a small, often blurred object that generic detectors miss. A common lightweight approach, used here, is to keep a generic detector and add a single-target motion model that rejects implausible candidates and bridges short gaps. Classical broadcast ball trackers also use motion directly: trajectory-based methods [13] generate many ball candidates and keep only the candidate chains whose motion is plausible for a ball. This work adds such a step for the ball in flight.

Related work at Amrita Vishwa Vidyapeetham, Bengaluru. YOLO detectors have been applied to online exam proctoring [15], classroom attentiveness monitoring [16] and a low-cost perception framework for autonomous driving [17]. Combining a deep detector with a tracker, Kumar *et al.* [14] detect and track multiple pedestrians, the closest task to player tracking, and Peri and Palaniswamy [18] track small, fast-moving animals with YOLOv8 and DeepSORT. Aarthi *et al.* [19] detect moving objects in surveillance video with a bottom-up approach, related to the motion stage used here for the ball in flight.

Detectors and metrics. YOLO detectors perform detection in a single network pass; this work uses the COCO-pretrained [5] YOLO11 nano model [4]. CLEAR-MOT [11] defines MOTA, which is dominated by detection errors. IDF1 [12] measures identity preservation. HOTA [10] balances detection and association explicitly and is the primary metric of SportsMOT. OWTA [20] removes the false-positive penalty for data in which not every object is labelled.

III. Problem Statement, Research Gap and Objectives

Problem. Given a monocular sports video, produce for every frame (i) a bounding box and a persistent identity for every player and (ii) the position of the ball, stating explicitly whether that position was detected in the frame or predicted by a motion model; and quantify detection quality, identity preservation, ball coverage and processing speed against baseline trackers under identical conditions.

Gap. Published sports trackers are usually evaluated with detectors fine-tuned on the benchmark and on GPUs; less is reported about how the classic trackers behave with a generic, lightweight, CPU-friendly detector, which is the realistic setting for a low-cost system. Players and the ball are also different tracking problems: players are many similar objects that need global assignment, whereas the ball is one tiny, frequently missed object, and benchmarks such as SportsMOT annotate players only. Finally, when a system interpolates the ball position, results frequently do not say which positions were observed and which were inferred. This motivates a hybrid design, a multi-object tracker for players and a dedicated single-target Kalman tracker for the ball, evaluated separately.

Objectives. (1) Detect players and the ball with a lightweight pretrained detector. (2) Track players with persistent identities using ByteTrack. (3) Track the ball with a Kalman filter that bridges short detection gaps without presenting a predicted position as a detection, and that follows the ball in flight when the detector loses it. (4) Export an annotated video, a per-frame table and a run summary. (5) Compare with SORT and OC-SORT on identical detections, frames and ground truth. (6) Evaluate on three public datasets with standard metrics. (7) Make every result reproducible.

IV. Proposed System

A. Overview

Fig. 1 shows the pipeline. Each frame is decoded and passed once through the detector, which reports the COCO classes *person* and *sports ball*. Inexpensive geometric filters remove implausible boxes. Player boxes go to the multi-object tracker and ball candidates to the ball tracker, which also receives motion-based ball flights computed from the frames. The tracked objects are drawn on the frame and streamed to a comma-separated file; a summary is written at the end. Players are many similar objects that need global assignment, whereas the ball is a single small object that is often missed, so the two are tracked by different algorithms: this combination is the proposed hybrid.

B. Detection and Filtering

YOLO11n (COCO weights, about 2.6 million parameters) is run at an input size of 1280 pixels with a non-maximum-suppression threshold of 0.5. Detections of *person* with confidence at least 0.15 and of *sports ball* with confidence at least 0.10 are kept. The thresholds are deliberately low because the trackers, not the detector threshold, decide what to trust. Player boxes must be between 12 pixels and 60% of the frame height tall with a height-to-width ratio between 0.8 and 6. Ball boxes must be 2–80 pixels in size and roughly square (aspect ratio at most 2). Sharing one network pass between the two classes halves the detection cost compared with two passes.

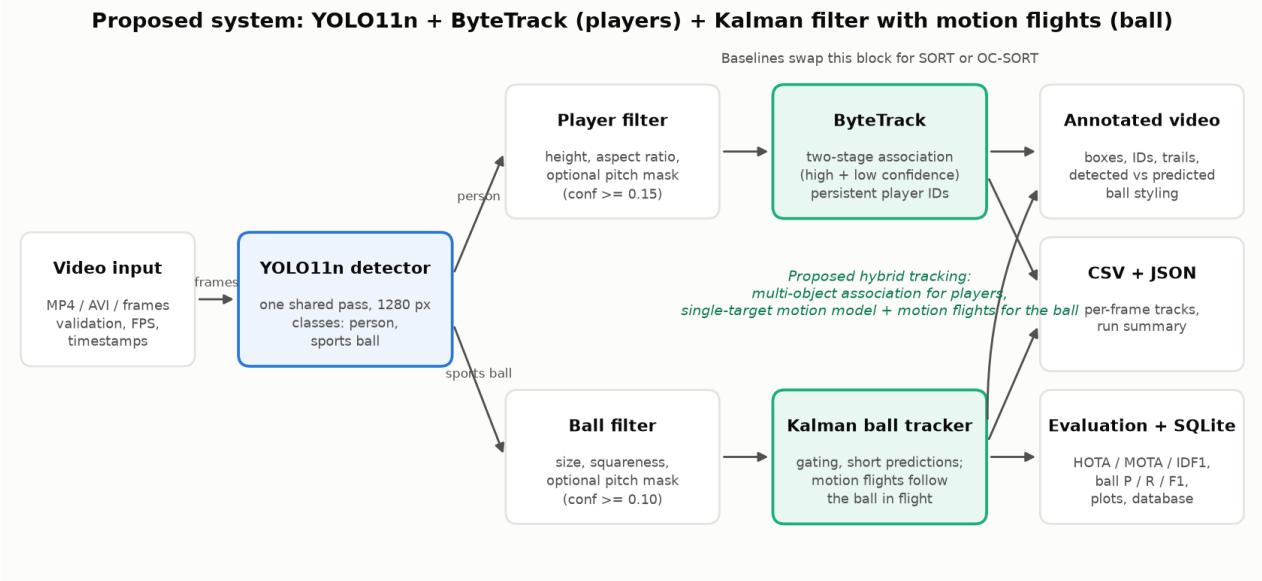


Fig. 1. Architecture of the proposed system. One YOLO11n pass per frame serves players and ball; ByteTrack tracks the players and a Kalman filter, extended with motion-based flights, tracks the ball. For the baselines only the player-tracker block is replaced by SORT or OC-SORT.

C. Player Tracking

SORT, OC-SORT and ByteTrack share one structure: each track is predicted with a Kalman filter, a cost matrix of intersection-over-union (IoU) between predicted boxes and detections is solved with the Hungarian algorithm, matched tracks are updated, unmatched detections start new tracks, and tracks that stay unmatched for a number of frames are deleted. They differ in how detections are used. SORT has a single association stage. OC-SORT uses only detections above a confidence threshold and adds an observation-centric re-update after occlusion and a direction-consistency term. ByteTrack first associates high-confidence detections and then uses the low-confidence ones to keep the remaining tracks alive.

ByteTrack is the proposed player tracker because, in sports, players are often partly occluded or motion-blurred, which lowers detector confidence without the player disappearing; ByteTrack is designed to keep such tracks alive instead of dropping them, which should reduce identity switches and fragmentation. It needs no appearance model, so it remains fast on a CPU and is not confused by identical kits.

D. Ball Tracking

Ball candidates come from the same detector pass. The ball tracker is a constant-velocity Kalman filter with state $\mathbf{x} = [x,\, y,\, \dot{x},\, \dot{y}]^\top$ and a time step of one frame:

$$\mathbf{x}_k = \mathbf{F}\mathbf{x}_{k-1} + \mathbf{w}_k, \qquad \mathbf{z}_k = \mathbf{H}\mathbf{x}_k + \mathbf{v}_k, \tag{1}$$

$$\mathbf{F} = \begin{bmatrix} 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}, \qquad \mathbf{H} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \end{bmatrix}, \tag{2}$$

where $\mathbf{w}_k$ models an unknown acceleration with a standard deviation of 4 pixels per frame squared and $\mathbf{v}_k$ a measurement error with a standard deviation of 4 pixels. Four rules turn the filter of (1)–(2) into a tracker.

Gating. A candidate is accepted only if its distance d_i to the predicted position is within a gate

$$g_k = \min\left(g_0 + m_k\,\Delta g,\; g_{\max}\right), \tag{3}$$

with $g_0 = 80$ pixels, $\Delta g = 20$ pixels, $g_{\max} = 160$ pixels and m_k the number of consecutive frames without an accepted detection. Among several candidates with confidence c_i the one that maximises

$$s_i = c_i \left(1 - \frac{d_i}{g_k}\right) \tag{4}$$

is taken. This rejects false balls elsewhere in the image.

Confirmation. A new track needs a detection with confidence at least 0.3 and is reported only after two detections, which suppresses one-frame false positives; weaker detections can only extend an existing track.

Prediction through gaps. With no accepted candidate, the filter coasts for up to 6 frames (0.2 s) and the output is labelled *predicted*. It does not coast right after a sudden jump of the measured position (a kick), because the new velocity is unknown.

Re-acquisition. A detection with confidence at least 0.5 outside the gate restarts the track after three missed frames, which handles long passes and camera cuts.

Predicted positions are kept distinct from end to end: a different marker and the text “ball (predicted)” in the video, a *predicted* status with an empty confidence in the exported table, separate counts in the run summary, and separate columns in every ball table of this paper.

E. Ball in Flight: Motion Flights

A struck ball is motion-blurred into a pale streak and the COCO detector almost never fires on it, while extrapolating the last velocity through a kick is wrong (Section VII-D). Three components follow the ball in flight.

Motion candidates. The two previous frames are aligned to the current one with a similarity transform estimated from ORB feature matches with RANSAC outlier rejection, which cancels camera pan and zoom. A pixel is moving if it differs from both aligned frames (three-frame differencing using only past frames, so the ghost at the ball's old position is suppressed). Bright, low-saturation moving pixels are grouped into components, and small, compact components are candidates. They are noisy by design: boots, socks and shirt numbers also move.

Flights. Candidates are linked frame to frame under a constant-velocity model. A chain counts as a flight while it has at least six points, a median speed of at least 8 pixels per frame and a smooth path, with a median change of velocity of at most $\max(4, 0.3\,\bar{v})$ pixels for median speed $\bar{v}$; this is rechecked every frame. Limbs swing back and forth and running players are slower, so they rarely qualify.

Adoption. A flight is used as the ball only if it was launched from the ball: its first velocity $\mathbf{v}_0$, run backwards from its first point $\mathbf{p}_0$ at frame t_0 to the frame t_s at which the detector last saw the ball at $\mathbf{b}$, must satisfy

$$\|\mathbf{p}_0 - \mathbf{v}_0\,(t_0 - t_s) - \mathbf{b}\| \le 35 \text{ pixels}. \tag{5}$$

Anchors are detector sightings only, so a wrong flight can never anchor the next one. A confident detection inside the gate prevents a new flight from being adopted, while a flight already being followed outranks single detections, which during a shot are often the goalkeeper's glove. When the followed flight ends abruptly (net, save, bounce) the ball is reported lost rather than extrapolated. A flight is confirmed only after six frames, so the pipeline holds its output eight frames and fills in the first frames of the flight. Flight positions are observations, reported as detected with the class name *ball (motion)* and no detector confidence.

V. Datasets

Three public datasets are used (Table I). Each has its own loader and annotations from different datasets are never mixed. The annotation format of every dataset was inspected on the downloaded files before its loader was written; UVY-Track, for example, is a multi-object-tracking text format with a class column, not a YOLO format. Only the required sequences were downloaded by reading the remote archives selectively (0.73 GB instead of about 11 GB). Table II lists the evaluated sequences: 4,850 frames in total (1,831 SportsMOT, 1,569 UVY-Track, 1,450 InFactory).

Tuning versus testing. One SportsMOT sequence (`v_2QhNRucNC7E_c017`) is the only data used to choose tracker hyper-parameters. It comes from a different source video than every test sequence and is never included in test averages. Nothing is tuned on UVY-Track or InFactory.

Two properties of the data matter for the results. First, SportsMOT labels only the players, so referees, coaches and perimeter staff detected by a generic *person* detector count as false positives. Second, UVY-Track labels only part of the people: in the fan-filmed penalty-kick clips `soccer_V01` and `soccer_V03` the four to six people involved in the kick have boxes, while dozens of visible photographers, staff and other players do not. Metrics that count false positives are therefore not valid error measures on UVY-Track. Its ball boxes are also loose, about twice the size of the detector boxes (20×26 against 12×11 pixels on `soccer_V03`), so an IoU of 0.5 is rarely reached even by a perfectly centred detection. InFactory has no ball identities, so only detection-type metrics are reported for it.

VI. Experimental Setup

A. Baselines and Fairness Protocol

Three player-tracking configurations are compared: YOLO11n + SORT, YOLO11n + OC-SORT, and the proposed YOLO11n + ByteTrack with the Kalman ball tracker. All three tracker implementations come from the open-source Roboflow `trackers` package (v2.6.1) and share one update interface, so no tracker was re-implemented. The protocol is as follows.

- The detector output is computed once per sequence and cached; every tracker replays exactly the same detections on the same frames and is scored against the same ground truth with the same evaluator.
- Hyper-parameters are chosen per tracker on the dev sequence with the same 48-point grid (track buffer in {30, 60} frames, frames before confirmation in {1, 3}, minimum IoU in {0.1, 0.2, 0.3}, confidence threshold in {0.3, 0.4, 0.5, 0.6}) and the same rule: highest HOTA, ties broken by IDF1. Table III lists the selected values.
- Results are also reported for one shared, untuned parameter set, so that the effect of tuning is visible.

For the ball, four outputs are compared on identical candidates: the most confident candidate of each frame with no temporal reasoning (*detection only*); the positions that the Kalman tracker reports from a detection, as an ablation; the Kalman tracker on detector candidates alone (detected and predicted positions, *YOLO only*), as an ablation without motion flights; and the full proposed tracker with motion flights, exactly as in the application. The Kalman part runs with its default configuration and is not tuned on any evaluated sequence. The motion-flight rules, however, were developed while inspecting the UVY-Track clip `soccer_V04`, and a first full run on InFactory exposed two failure modes (a camera cut and a fast object passing the ball) that led to two further rules, so the ball results of the proposed method are not from fully held-out data. A further clip, `soccer_V02`, never used for any decision, was checked qualitatively.

B. Metrics

Player tracking is scored with the evaluation module of the `trackers` package, a re-implementation designed to

TABLE I
DATASETS AND HOW THEY ARE USED

	SportsMOT [6]	UVY-Track [7]	InFactory Soccer Ball Tracking [8]
Footage	Professional broadcast (football, basketball, volleyball)	User-generated online sports videos, seven sports	Professional soccer broadcast
Annotation format	MOT Challenge text	MOT-style text with a class column	YOLO text, one file per frame
Annotated objects	Players on the pitch only	Player, goalkeeper, referee, sports ball, goal, camera	Ball only, plus frames marked "ball not visible"
Persistent identities	Players	People and ball	None
Role in this work	Player tracking (tuning and test)	Player tracking and ball detection/tracking (test)	Ball detection (test)
Subset used	4 football sequences of the validation split	3 soccer sequences	10 clips (884 frames with a visible ball, 566 without)

TABLE II
EVALUATED SEQUENCES

Dataset	Sequence	Role	Frames	Tracks	Ball fr.
SportsMOT	`v_2QhNRucNC7E_c017`	dev	450	21	–
SportsMOT	`v_dw7LOz17Omg_c067`	test	450	22	–
SportsMOT	`v_G-vNjfx1GGc_c008`	test	480	20	–
SportsMOT	`v_i2_L4qquVg0_c007`	test	451	21	–
UVY-Track	`soccer_V04`	test	219	19	168
UVY-Track	`soccer_V01`	test	750	6	750
UVY-Track	`soccer_V03`	test	600	5	600
InFactory	10 clips	test	1,450	–	884

TABLE III
TRACKER PARAMETERS (N/A: PARAMETER NOT USED BY THE TRACKER)

Parameter set	Buffer (frames)	Confirm (frames)	Min. IoU	New-track threshold	High-conf. threshold
Shared default	30	3	0.1	0.5	0.5
SORT, tuned	30	1	0.3	0.6	n/a
OC-SORT, tuned	60	1	0.3	n/a	0.3
ByteTrack (proposed), tuned	30	1	0.2	0.5	0.5

match TrackEval, the reference evaluator of MOTChallenge and SportsMOT. HOTA [10] is the primary metric,

$$\mathrm{HOTA}_\alpha = \sqrt{\mathrm{DetA}_\alpha \cdot \mathrm{AssA}_\alpha}, \tag{6}$$

averaged over localisation thresholds α from 0.05 to 0.95, where DetA and AssA are the detection and association accuracies. MOTA [11] and IDF1 [12] are

$$\mathrm{MOTA} = 1 - \frac{\mathrm{FN} + \mathrm{FP} + \mathrm{IDSW}}{\mathrm{GT}}, \tag{7}$$

$$\mathrm{IDF1} = \frac{2\,\mathrm{IDTP}}{2\,\mathrm{IDTP} + \mathrm{IDFP} + \mathrm{IDFN}}, \tag{8}$$

computed at an IoU threshold of 0.5. Also reported are identity switches (IDSW), false positives (FP), false negatives (FN), fragmentations (Frag) and mostly tracked targets (MT). For the partially labelled UVY-Track data the false-positive penalty is removed by using [20]

$$\mathrm{OWTA}_\alpha = \sqrt{\mathrm{DetRe}_\alpha \cdot \mathrm{AssA}_\alpha}, \tag{9}$$

together with the detection recall DetRe.

Ball and detector metrics are computed with an evaluator written for this project. It is unit-tested on hand-computed cases but is not an official benchmark tool. A method outputs at most one ball position per frame. An output is correct under the *distance* criterion if its centre lies within one ball diameter of the true centre, and under the *IoU* criterion if its IoU with a ground-truth box is at least 0.5. Precision, recall (the visible-ball recall) and F1 follow from the per-frame outcomes. AP@0.5 is the average precision of all ball candidates ranked by confidence; with one class it equals mAP@0.5. Coverage is the share of labelled frames with any output. Continuity is described by the number and mean length of correctly tracked segments inside runs of visible-ball frames. Ball identity switches are reported only for UVY-Track, which has ball identities.

C. Implementation and Hardware

All experiments ran on an Apple M1 CPU (8 physical / 8 logical cores), 8.0 GB RAM, no GPU used (device: cpu), macOS, Python 3.14.6, PyTorch 2.15.0 with 7 threads, Ultralytics 8.4.198, `trackers` 2.6.1, `supervision` 0.31.3 and OpenCV 5.0.0. The system has 121 automated tests, including an end-to-end test of the full pipeline. All results are reproduced by one command in the project repository[1]; the evaluated sequences and frame ranges are fixed in a manifest file. The numbers are computed on small subsets with a generic detector and are not comparable with leaderboard entries.

VII. RESULTS

A. Player Detection

Table IV compares the detector output with the ground-truth player boxes before any tracking. Recall is above 92% on the SportsMOT test sequences. Precision is limited by real people who are not annotated, and on UVY-Track it cannot be interpreted at all for the reason given in Section V.

B. Player Tracking on SportsMOT

Table V and Fig. 2 give the results on the three SportsMOT test sequences (1,381 frames, 16,360 ground-truth boxes).

[1]https://github.com/YaswanthKumar814/Multi-Object-Tracking-for-Sports-Analytics-

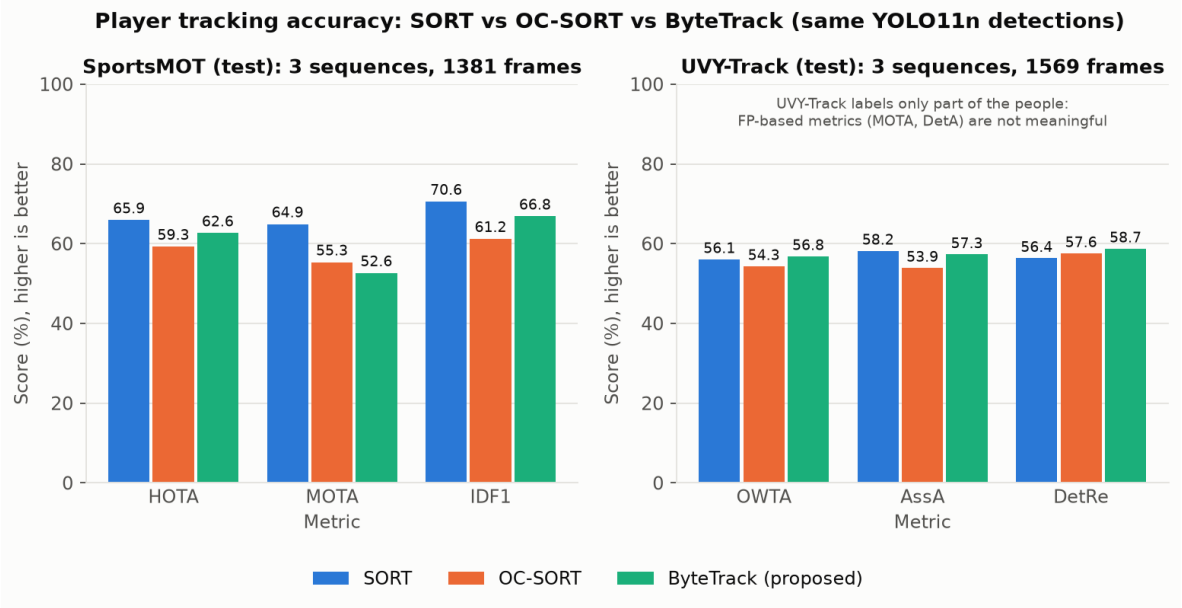


Fig. 2. Tracking accuracy of the three trackers with tuned parameters. Left: HOTA, MOTA and IDF1 on SportsMOT. Right: OWTA, AssA and detection recall on the partially labelled UVY-Track.

TABLE IV
PLAYER DETECTION WITH YOLO11N (CONFIDENCE ≥ 0.5, IOU ≥ 0.5)

Dataset	Role	Seq.	Prec.	Rec.	F1	AP@0.5
SportsMOT	dev	1	92.4	89.3	90.8	90.1
SportsMOT	test	3	81.3	92.2	86.4	88.9
UVY-Track	test	3	49.6†	70.8	58.3†	57.4†

†Not meaningful: only part of the people are labelled.

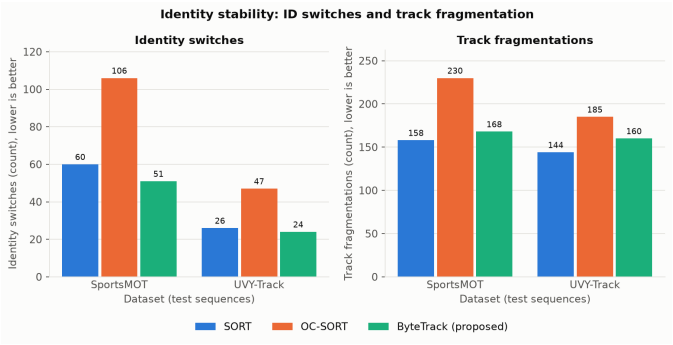


Fig. 3. Identity switches and track fragmentations on the test sequences of both datasets (lower is better).

Identity stability. ByteTrack has the fewest identity switches (51), the fewest missed boxes (670 false negatives, recall 95.9%) and the most mostly tracked players. OC-SORT has about twice as many switches (106) and the most fragmented tracks (Fig. 3).

Overall accuracy. Tuned SORT has the best HOTA (65.9), MOTA (64.9) and IDF1 (70.6). ByteTrack is second on HOTA (62.6) and IDF1 (66.8). Its association accuracy is within 0.9 points of SORT (AssA 66.1 against 67.0); the gap comes from detection accuracy, because ByteTrack outputs 2,131 more false-positive boxes.

Explanation. The second association stage of ByteTrack keeps tracks alive with low-confidence boxes. With a detector trained on the benchmark those boxes are occluded players. With a generic detector on SportsMOT, many of them are unannotated people at the edge of the pitch, so the mechanism that lowers the number of misses also raises the number of false positives: 1,912 of the 2,131 extra false positives come from the single sequence with staff along the perimeter boards. In addition, the SORT implementation used here already allows any detection, including a low-confidence one, to match an existing track, so the practical difference between the two trackers is smaller than in the original papers; the tuned SORT settings (new tracks only above a confidence of 0.6, IoU at least 0.3) simply start fewer tracks on bystanders.

Effect of tuning. With one shared, untuned parameter set (lower half of Table V) the three trackers are within 0.8 HOTA points of each other. ByteTrack again has the fewest identity switches (43). OC-SORT then has the best MOTA (71.1) because it discards every box below a confidence of 0.5, which removes many false positives but costs about 8 points of recall. On the dev sequence itself ByteTrack was best on HOTA, MOTA, IDF1 and identity switches, but this is not an independent result. The selected parameters did not transfer equally well: the dev sequence contains few unannotated people, so a low confidence threshold looked best for OC-SORT there and then hurt it on the test sequences. Fig. 4 shows the sensitivity of each tracker to its parameters. Tuning on a single sequence is a limitation of this study.

TABLE V
PLAYER TRACKING ON THE SPORTSMOT TEST SEQUENCES WITH IDENTICAL YOLO11N DETECTIONS. BEST VALUE PER COLUMN AND PARAMETER SET IN BOLD (LOWER IS BETTER FOR IDSW, FP, FN AND FRAG)

Method	HOTA	DetA	AssA	MOTA	IDF1	IDSW	FP	FN	Frag	MT	Prec.	Rec.
Parameters tuned per tracker on the dev sequence												
YOLO11n + SORT	**65.9**	**64.8**	**67.0**	**64.9**	**70.6**	60	**4,898**	784	**158**	59	**76.1**	95.2
YOLO11n + OC-SORT	59.3	60.6	58.0	55.3	61.2	106	6,382	817	230	61	70.9	95.0
YOLO11n + ByteTrack (proposed)	62.6	59.4	66.1	52.6	66.8	**51**	7,029	**670**	168	**62**	69.1	**95.9**
One shared, untuned parameter set (project defaults)												
YOLO11n + SORT	**64.1**	59.5	**69.0**	52.9	68.6	44	6,975	**687**	**174**	**61**	69.2	**95.8**
YOLO11n + OC-SORT	63.5	**67.4**	59.9	**71.1**	**69.4**	61	**2,629**	2,030	259	52	**84.5**	87.6
YOLO11n + ByteTrack (proposed)	63.3	59.4	67.4	52.9	67.8	**43**	6,963	696	**174**	**61**	69.2	95.7

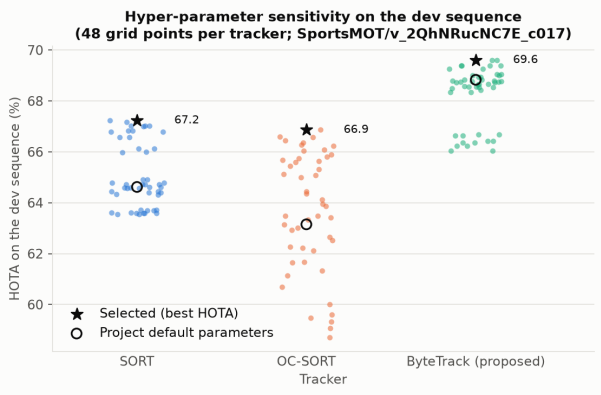


Fig. 4. HOTA on the dev sequence for the 48 grid points of each tracker. Star: selected setting. Circle: shared default setting.

TABLE VI
PLAYER TRACKING ON THE UVY-TRACK TEST SEQUENCES (1,569 FRAMES, TUNED PARAMETERS). ONLY METRICS THAT DO NOT COUNT FALSE POSITIVES ARE SHOWN

Tracker	OWTA	DetRe	AssA	IDSW	Frag	MT	FN
SORT	56.1	56.4	**58.2**	26	**144**	**16**	2,524
OC-SORT	54.3	57.6	53.9	47	185	**16**	2,363
ByteTrack (proposed)	**56.8**	**58.7**	57.3	**24**	160	**16**	**2,232**

C. Player Tracking on UVY-Track

Table VI reports the metrics that remain valid when only part of the people are labelled. On these the trackers are close. ByteTrack has the best OWTA (56.8), the highest detection recall (58.7), the fewest identity switches (24) and the fewest missed boxes; SORT has a marginally higher AssA and the fewest fragmentations; OC-SORT has almost twice as many identity switches (47). For completeness, MOTA on this dataset is negative for every tracker (between -32 and -93), which reflects the unlabelled people and not tracker errors. The low-resolution clip `soccer_V01` (636×360 pixels, people about 35 pixels tall) is difficult for all three trackers, with an OWTA of about 30.

D. Ball Detection and Tracking

Table VII summarises the ball results on 13 sequences (3,019 labelled frames, 2,402 with a visible ball); Fig. 5 shows precision, recall and F1.

The detector is the bottleneck. The COCO *sports ball* class finds the ball in only about a third of the frames in which it is visible (recall 34.8%; AP@0.5 of 16.3% over all candidates). When it does respond to the ball it is accurate: correct outputs lie about 2 pixels from the true centre.

Effect of the Kalman tracker. Over detection alone, precision rises from 66.4% to 71.1%, recall from 34.8% to 35.8% and F1 from 45.7% to 47.6% under the distance criterion, and F1 improves on both datasets. The tracker is deliberately conservative (a confidence of at least 0.3 to start a track, a capped gate, at most six coasted frames and none right after a kick). A looser earlier version reached a higher recall (43.3%) and F1 (49.3%), but only 42.6% of its predicted positions were correct; the current version trades some recall for positions that can be trusted. The ball track also becomes far more continuous: with the proposed tracker 175 short correct segments merge into 44 and their mean length grows from 4.8 to 19.9 frames. It recovers 249 visible-ball frames that detection alone misses and loses 212 that detection alone had (the first detection of each track, which is held back for confirmation, weak detections that may not start a track, and detections outside the gate). On UVY-Track, the only dataset with ball identities, the track identity changes four times.

Effect of motion flights. On top of the Kalman tracker, recall rises from 35.8% to 36.4% and F1 from 47.6% to 47.9%. Flights fire only where the ball is struck and visible in flight: on every InFactory clip and on two of the three UVY-Track clips the output is identical to the Kalman tracker. On the free kick of `soccer_V04` the Kalman tracker loses the ball when it is struck (recall 91.1%), whereas with motion flights the shot is followed from the boot into the net (recall 99.4%, Fig. 6). Precision there falls from 97.5% to 84.8%, but all 24 outputs counted as false lie in frames in which the annotators marked no ball (the most blurred part of the flight and the ball inside the net); inspection shows these positions on the ball.

Table VIII separates detected and predicted outputs. Detected positions accepted by the tracker are more precise than raw detections (73.2% against 66.4%), because gating and confirmation reject isolated false balls. Predicted positions are correct 60.1% of the time: they are right when the ball rolls or is briefly occluded and wrong when it changes direction during

TABLE VII

BALL RESULTS: DETECTION ONLY VERSUS THE KALMAN BALL TRACKER WITH AND WITHOUT MOTION FLIGHTS ON IDENTICAL CANDIDATES. AP@0.5 IS COMPUTED OVER ALL DETECTOR CANDIDATES. COVERAGE IS THE SHARE OF LABELLED FRAMES WITH AN OUTPUT. FRAMES: DETECTED, PREDICTED, NO OUTPUT. SEGMENTS: NUMBER AND MEAN LENGTH IN FRAMES

		Distance criterion (%)			IoU ≥ 0.5 (%)			AP@0.5	Cover.	Frames			Correct segments	
Data	Method	Prec.	Rec.	F1	Prec.	Rec.	F1	(%)	(%)	Det.	Pred.	None	No.	Mean len.
UVY-Track	Detection only	62.6	35.1	45.0	24.8	13.9	17.8	12.0	54.3	852	0	717	102	5.2
UVY-Track	Kalman, detected only	67.0	28.1	39.6	34.2	14.4	20.2	–	40.6	637	0	932	43	9.9
UVY-Track	Kalman (YOLO only)	65.6	35.2	45.8	27.8	14.9	19.4	–	51.9	637	177	755	10	53.4
UVY-Track	Kalman + motion (proposed)	64.2	36.1	46.2	28.1	15.8	20.2	–	54.4	672	182	715	11	49.8
InFactory	Detection only	74.5	34.4	47.1	66.9	30.9	42.3	26.0	28.1	408	0	1,042	73	4.2
InFactory	Kalman, detected only	92.5	27.8	42.8	80.1	24.1	37.0	–	18.3	266	0	1,184	50	4.9
InFactory	Kalman (YOLO only)	82.5	36.9	51.0	63.5	28.4	39.2	–	27.2	266	129	1,055	33	9.9
InFactory	Kalman + motion (proposed)	82.5	36.9	51.0	63.5	28.4	39.2	–	27.2	266	129	1,055	33	9.9
All (13 seq.)	Detection only	66.4	34.8	45.7	38.4	20.1	26.4	16.3	41.7	1,260	0	1,759	175	4.8
All (13 seq.)	Kalman, detected only	74.5	28.0	40.7	47.7	17.9	26.1	–	29.9	903	0	2,116	93	7.2
All (13 seq.)	Kalman (YOLO only)	71.1	35.8	47.6	39.5	19.9	26.4	–	40.0	903	306	1,810	43	20.0
All (13 seq.)	Kalman + motion (proposed)	70.0	36.4	47.9	39.3	20.4	26.9	–	41.4	938	311	1,770	44	19.9

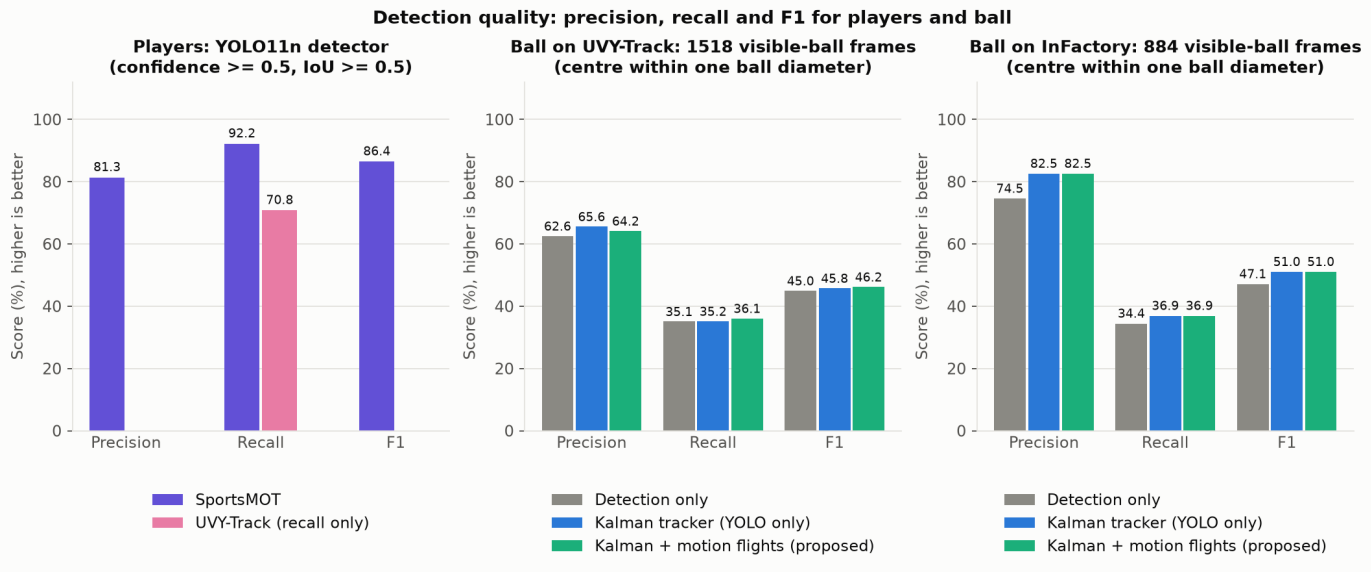


Fig. 5. Precision, recall and F1. Left: player detection. Centre and right: the ball on UVY-Track and on InFactory: detection only, the Kalman tracker on detector candidates alone, and the proposed tracker with motion flights (distance criterion).

TABLE VIII

PROPOSED BALL TRACKER: PRECISION OF DETECTED AND OF PREDICTED POSITIONS (DISTANCE CRITERION)

	Detected		Predicted	
Data	Frames	Prec. (%)	Frames	Prec. (%)
UVY-Track	672	65.6	182	58.8
InFactory	266	92.5	129	62.0
All	938	73.2	311	60.1

the gap, because the filter extrapolates the last velocity. This is the reason for reporting predicted positions separately.

Strict localisation. Under the IoU criterion F1 barely changes (26.4% to 26.9%), because a predicted box is rarely within an IoU of 0.5 of a ball that is about 15 pixels wide. On UVY-Track the IoU-based numbers are low for every method because of the loose ground-truth boxes described in Section V; the distance criterion is the meaningful one there.

Per-sequence behaviour is uneven. Against detection only, F1 improves on five sequences (for example from 45.6% to 78.7% on `71125020`), stays at zero on three clips in which the detector never finds the ball, and decreases on five (for example from 51.5% to 47.5% on `soccer_V03`, and on `7d503a7a`, where the track follows something other than the ball). Fig. 7 shows the share of detected, predicted and lost frames per sequence.

E. Detector Comparison

As a secondary experiment, YOLOv8n and YOLO11n were compared with identical settings on the first 200 frames of one SportsMOT sequence (players) and on two InFactory clips

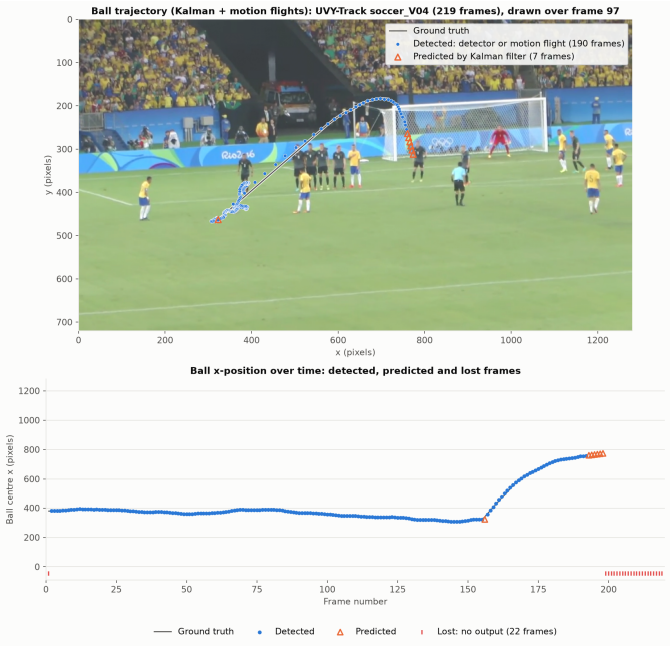


Fig. 6. Ball trajectory of the proposed tracker on the free kick of `soccer_V04` with ground truth. Dots: detected, by the detector or by a motion flight. Triangles: predicted by the Kalman filter. Marks on the time axis: frames without output. The arc is the shot, followed by a motion flight; the ground truth has a gap during the most blurred part of the flight.

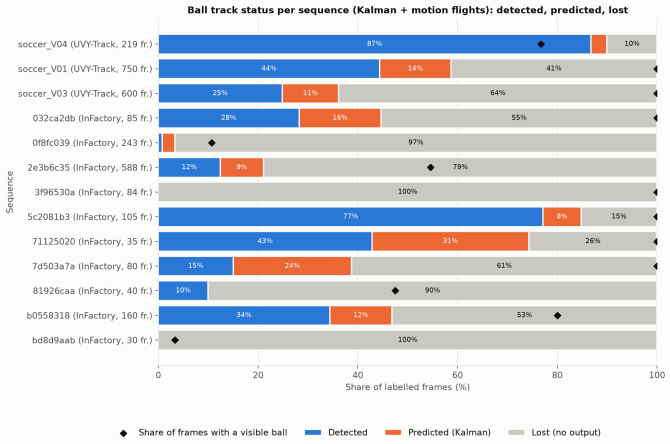


Fig. 7. Ball track status per sequence with the proposed tracker. Diamonds mark the share of frames in which the ball is visible.

(ball). Speed was measured by alternating the two models frame by frame so that both saw the same machine state (Table IX). On this subset YOLO11n is clearly better for players (F1 of 81.2% against 71.7%), mainly through higher precision, and about 5% faster. For the ball the two models are practically equal. The two ball clips are among the easier ones, so their ball scores are above the averages of Table VII. This supports the choice of YOLO11n but is a small-sample comparison.

TABLE IX
YOLOv8n Versus YOLO11n on the Same Frames

Model	Target	Frames	Prec.	Rec.	F1	AP@0.5	ms/fr.
YOLOv8n	player	200	60.9	87.1	71.7	71.7	112
YOLOv8n	ball	265	83.0	67.0	74.1	51.0	112
YOLO11n	player	200	74.5	89.2	81.2	79.0	106
YOLO11n	ball	265	88.8	64.8	74.9	51.2	106

TABLE X
Median Processing Time per Frame

Stage	Time (ms)	Share (%)
Detection	112.63	64.0
Ball motion: camera compensation + flight linking	52.83	30.0
Video encoding	4.45	2.5
Frame reading	4.32	2.5
Drawing annotations	1.11	0.6
Player filter + ByteTrack	0.71	0.4
Ball filter + Kalman tracker	0.03	0.0
Total (proposed pipeline)	**176.1**	100.0
Player filter + SORT (baseline)	0.61	–
Player filter + OC-SORT (baseline)	1.59	–

F. Runtime

Table X and Fig. 8 give the median time per frame of every stage, measured over 150 frames of a 1280×720 sequence on an idle machine on mains power. The detector takes 64% of the time and the ball motion stage (feature-based frame alignment, differencing and flight linking) another 30%. The player trackers cost about a millisecond, so the choice of tracker changes throughput by under 1%: the full pipeline runs at 5.68 frames per second with ByteTrack, 5.68 with SORT and 5.65 with OC-SORT. SORT is the cheapest tracker, ByteTrack costs slightly more for its second association stage, and OC-SORT is two to three times slower than both. The system is therefore an offline analysis tool on this hardware; disabling motion flights removes the motion stage. Power state matters: on battery the processor is throttled and every timing changes by a large factor, so the power source is recorded with the results.

G. Qualitative Results

The application was run end to end on the UVY-Track clip `soccer_V04` (241 frames, a free kick and a goal). It processed the clip in 46 seconds (5.2 frames per second), produced 33 player track identities with a median track length of 209 frames, and reported the ball in 197 frames, of which 190 were detected (by the detector or a motion flight) and 7 predicted; the shot is followed into the net. On a further, previously unseen fan-filmed penalty shoot-out (`soccer_V02`, 1,085 frames) a penalty is likewise followed from the spot past the diving goalkeeper into the net. Fig. 9 shows a frame of the annotated output. Figs. 10 and 11 show player trajectories and a position heatmap for a SportsMOT sequence; both are in image coordinates, so camera motion is not compensated.

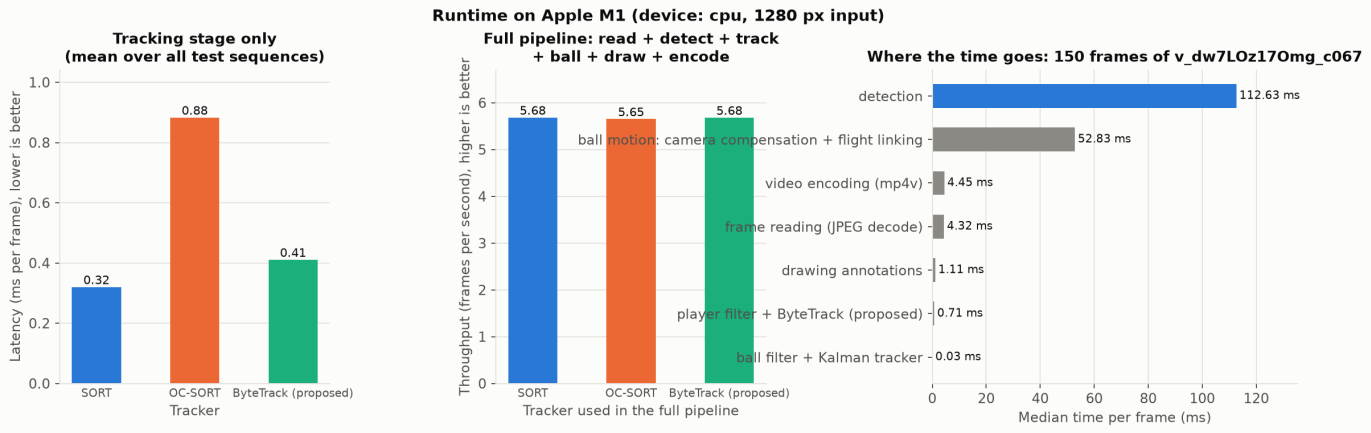


Fig. 8. Runtime. Left: latency of the tracking stage alone. Centre: throughput of the full pipeline. Right: time per stage of the proposed system.

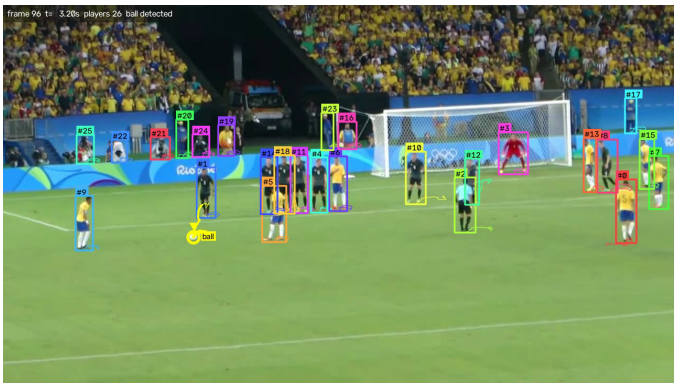


Fig. 9. A frame of the annotated output video: player boxes with identities and trails, the ball marker with its recent path, and the status line.

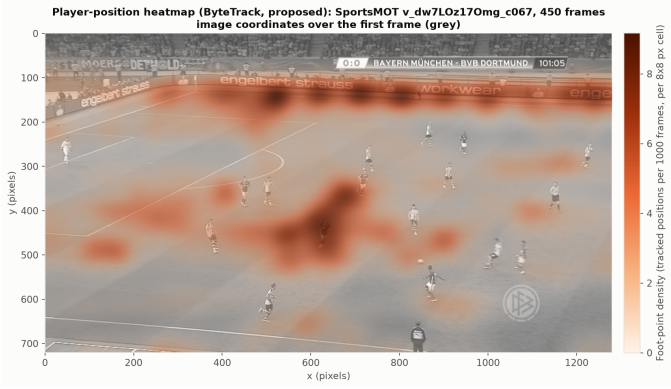


Fig. 11. Player-position heatmap over the same sequence (image coordinates).

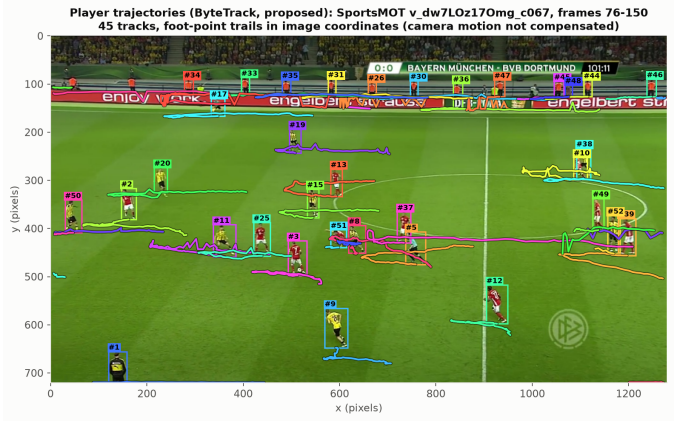


Fig. 10. Player trajectories from ByteTrack over 75 frames of a SportsMOT test sequence (foot points, image coordinates). The tracks along the top are perimeter staff, who are detected and tracked but are not annotated in the benchmark.

VIII. Limitations

- *Generic detector.* The *person* class also responds to referees, coaches, staff and spectators, and the *sports ball* class misses most small broadcast balls. This limits every number in this paper more than the choice of tracker does.
- *The proposed tracker is not uniformly best.* ByteTrack gives the most stable identities and the highest recall, but tuned SORT has a higher HOTA, MOTA and IDF1 on SportsMOT in this setting.
- *Small evaluation.* Six tracking test sequences, thirteen ball sequences and a single tuning sequence; no confidence intervals are given and the results are not comparable with published leaderboards.
- *Partial labels.* No single dataset evaluates the whole system: SportsMOT has no ball labels, UVY-Track labels only part of the people, and InFactory has neither player labels nor ball identities.
- *Ball predictions.* A constant-velocity model cannot follow a direction change; 40% of the predicted positions are wrong, and the tracker sometimes follows a false positive or a second real ball beside the pitch.
- *Motion flights are narrow.* They help only when the ball is struck and visible in flight, must start where the detector last saw the ball, and are switched off when the global camera model fails (cuts, dissolves, fast zoom). Their rules were developed on UVY-Track and refined after a first InFactory run, because no separate ball tuning set

was available.

- *Scope.* There is no camera-motion compensation for the players, re-identification, team assignment or pitch calibration, and throughput is about six frames per second on an Apple M1 processor.

IX. Future Work

The most valuable next step is a soccer-specific detector with separate classes for players, referees, goalkeepers and the ball, which addresses the dominant error source for both players and ball. Tiled high-resolution inference for the ball is already supported by the application but has not been evaluated. The ball motion model can be improved with a gravity-aware model for flights and offline forward-backward smoothing, and motion flights should be evaluated on a dedicated, fully labelled ball tuning and test split. Camera-motion compensation and appearance-based re-identification would steady trajectories and identities. Further steps are team assignment from jersey colour, pitch calibration for positions in metres, accelerated inference for real-time throughput, and evaluation on the complete validation splits with several tuning sequences.

X. Conclusion

A complete player-and-ball tracking system was built, tested and evaluated on three public datasets with a reproducible protocol. With identical detections, ByteTrack produced the fewest identity switches and the highest recall on both tracking datasets and the best OWTA on UVY-Track, while tuned SORT obtained the best HOTA, MOTA and IDF1 on SportsMOT; OC-SORT was the weakest in association and the slowest. The dedicated Kalman ball tracker improved ball precision, F1 and above all the continuity of the ball track over detection alone, and motion-based flights let it follow a struck ball that the detector loses without changing the result on clips without such a shot. Its predicted positions are still wrong 40% of the time, which justifies labelling them separately. Across all experiments the generic lightweight detector, not the tracker, limits both accuracy and speed, so a soccer-specific detector is the most promising improvement.

References

[1] A. Bewley, Z. Ge, L. Ott, F. Ramos, and B. Upcroft, "Simple online and realtime tracking," in *Proc. IEEE Int. Conf. Image Process. (ICIP)*, 2016, pp. 3464–3468.

[2] Y. Zhang, P. Sun, Y. Jiang, D. Yu, F. Weng, Z. Yuan, P. Luo, W. Liu, and X. Wang, "ByteTrack: Multi-object tracking by associating every detection box," in *Proc. Eur. Conf. Comput. Vis. (ECCV)*, 2022.

[3] J. Cao, J. Pang, X. Weng, R. Khirodkar, and K. Kitani, "Observation-centric SORT: Rethinking SORT for robust multi-object tracking," in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2023.

[4] G. Jocher and J. Qiu, "Ultralytics YOLO11," software, 2024. [Online]. Available: https://github.com/ultralytics/ultralytics

[5] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollár, and C. L. Zitnick, "Microsoft COCO: Common objects in context," in *Proc. Eur. Conf. Comput. Vis. (ECCV)*, 2014, pp. 740–755.

[6] Y. Cui, C. Zeng, X. Zhao, Y. Yang, G. Wu, and L. Wang, "SportsMOT: A large multi-object tracking dataset in multiple sports scenes," in *Proc. IEEE/CVF Int. Conf. Comput. Vis. (ICCV)*, 2023.

[7] E. Alencar and R. de Freitas, "UVY: A dataset of sport user-generated videos from YouTube for multi-object tracking evaluation," Zenodo, 2026. [Online]. Available: https://zenodo.org/records/21303900

[8] V. Constantinou, M. Iranmanesh, and J. Kanalakis, "Soccer ball tracking dataset," Infactory, 2026. [Online]. Available: https://huggingface.co/datasets/infactory-ai/ball-tracking

[9] R. E. Kalman, "A new approach to linear filtering and prediction problems," *J. Basic Eng.*, vol. 82, no. 1, pp. 35–45, 1960.

[10] J. Luiten, A. Ošep, P. Dendorfer, P. Torr, A. Geiger, L. Leal-Taixé, and B. Leibe, "HOTA: A higher order metric for evaluating multi-object tracking," *Int. J. Comput. Vis.*, vol. 129, 2021.

[11] K. Bernardin and R. Stiefelhagen, "Evaluating multiple object tracking performance: The CLEAR MOT metrics," *EURASIP J. Image Video Process.*, 2008.

[12] E. Ristani, F. Solera, R. Zou, R. Cucchiara, and C. Tomasi, "Performance measures and a data set for multi-target, multi-camera tracking," in *Proc. Eur. Conf. Comput. Vis. Workshops (ECCVW)*, 2016.

[13] X. Yu, H. W. Leong, C. Xu, and Q. Tian, "Trajectory-based ball detection and tracking in broadcast soccer video," *IEEE Trans. Multimedia*, vol. 8, no. 6, 2006.

[14] A. Kumar, T. Singh, and P. Duraisamy, "Detection and tracking of multiple pedestrians using deep learning," in *Proc. 14th Int. Conf. Comput. Commun. Netw. Technol. (ICCCNT)*, 2023, pp. 1–6.

[15] T. Singh, R. R. Nair, T. Babu, and P. Duraisamy, "Enhancing academic integrity in online assessments: Introducing an effective online exam proctoring model using YOLO," *Procedia Comput. Sci.*, vol. 235, pp. 1399–1408, 2024.

[16] R. R. Nair, T. Babu, T. Singh, P. Duraisamy, and V. Mehra, "Class room student attentiveness model based on YOLO," in *Proc. 14th Int. Conf. Comput. Commun. Netw. Technol. (ICCCNT)*, 2023.

[17] V. Gopi Kiran, T. Hemanth Babu, N. A. Vidula, and T. Singh, "A low-cost vision-based framework for autonomous driving using YOLO, MiDaS, and stereo vision," in *IFIP Advances in Information and Communication Technology*, Springer, 2025, doi: 10.1007/978-3-031-98360-3_4.

[18] S. D. B. Peri and S. Palaniswamy, "A novel approach to detect and track small animals using YOLOv8 and DeepSORT," in *Proc. 4th IEEE Global Conf. Advancement Technol. (GCAT)*, 2023, pp. 1–6, doi: 10.1109/GCAT59970.2023.10353296.

[19] R. Aarthi, J. Amudha, K. Boomika, and A. Varrier, "Detection of moving objects in surveillance video by integrating bottom-up approach with knowledge base," *Procedia Comput. Sci.*, vol. 78, pp. 160–164, 2016.

[20] Y. Liu, I. E. Zulfikar, J. Luiten, A. Dave, D. Ramanan, B. Leibe, A. Ošep, and L. Leal-Taixé, "Opening up open world tracking," in *Proc. IEEE/CVF Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2022.